

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing, svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

In [2]:

```
1 df=pd.read_csv(r"C:\Users\teppa\Downloads\USA_Housing.csv")
2 df
```

Out[2]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Fer 674\nLaurab
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Suite 079 Kathleer
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Eli Stravenue\nDanik WI 0
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nF
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond AE
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams AP 3015:
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 925 8489\nAPO AA
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy C Suite 076\nJoshu V
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nF
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George I Apt. 509\nEas

5000 rows × 7 columns



In [3]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   Avg. Area Income                       5000 non-null   float64
 1   Avg. Area House Age                    5000 non-null   float64
 2   Avg. Area Number of Rooms              5000 non-null   float64
 3   Avg. Area Number of Bedrooms           5000 non-null   float64
 4   Area Population                         5000 non-null   float64
 5   Price                                  5000 non-null   float64
 6   Address                                5000 non-null   object 
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [4]:

```
1 df.describe()
```

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06


In [5]:

```
1 df.columns
```

Out[5]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Addresses'], dtype='object')



```
In [11]: 1 sns.displot(df['Price'])
```

```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x14cd96e0d00>
```



```
In [10]: 1 sns.displot(df['Area Population'])
```


```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x14cd96c31c0>
```



```
In [29]: 1 Housedf=df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of  
2           'Avg. Area Number of Bedrooms', 'Area Population', 'Price']]
```

```
In [30]: 1 sns.heatmap(Housedf.corr())
```

Out[30]: <Axes: >



```
In [32]: 1 x=Housedf[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
2             'Avg. Area Number of Bedrooms', 'Area Population']]  
3 y=df['Price']
```

```
In [35]: 1 from sklearn.model_selection import train_test_split  
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
```

```
In [36]: 1 from sklearn.linear_model import LinearRegression  
2 lm=LinearRegression()  
3 lm.fit(x_train,y_train)  
4
```

Out[36]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [37]: 1 print(lm.intercept_)
```

-2641372.6673014304


```
In [50]: 1 coeff_df=pd.DataFrame(lm.coef_,x.columns,columns=['coefficient'])
2 coeff_df
```

```
Out[50]:
```


	coefficient
Avg. Area Income	21.617635
Avg. Area House Age	165221.119872
Avg. Area Number of Rooms	121405.376596
Avg. Area Number of Bedrooms	1318.718783
Area Population	15.225196

```
In [42]: 1 predictions=lm.predict(x_test)
2 plt.scatter(y_test,predictions)
```

```
Out[42]: <matplotlib.collections.PathCollection at 0x14cd9e13700>
```



```
In [44]: 1 sns.displot((y_test-predictions),bins=50);
```



```
In [46]: 1 from sklearn import metrics
2 print('MAE:',metrics.mean_absolute_error(y_test,predictions))
3 print('MSE:',metrics.mean_squared_error(y_test,predictions))
4 print('MSE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

```
MAE: 81257.55795855941
MSE: 10169125565.897606
MSE: 100842.08231635048
```

```
In [ ]: 1
```